

The background is a dark gray architectural floor plan. It features various geometric shapes representing rooms, corridors, and structural elements. Faint white lines and text are visible throughout the plan, including room numbers like '108', '110', and '109', and technical specifications such as '2 935.8347', '4 500', '13 622.9226', '5 023.4567', and '6090 / C110 TOIMISTO 440.0 m²'.

CPU Caches: Trust but Verify

Tommi Jalkanen - Kodan

Rust Finland

- First meetup in February 3, 2017
- 15 meetups organized, 277 members on meetup.com and 64 on matrix
- Chat with others: #rust-finland:matrix.org
- Website: <https://www.rust-finland.org>
- Github organization: <https://github.com/rust-finland>

Motivation

- Sharing knowledge
- Bad takes on the Internet
- Time + money + environment

Bottleneck for Computing

Bottleneck for Computing

- Speed of light

Bottleneck for Computing

- Speed of light
- 2 GHz CPU can do ~ 2 instructions/ns

Bottleneck for Computing

- Speed of light
- 2 GHz CPU can do ~ 2 instructions/ns
- Light travels ~ 0.3 m/ns

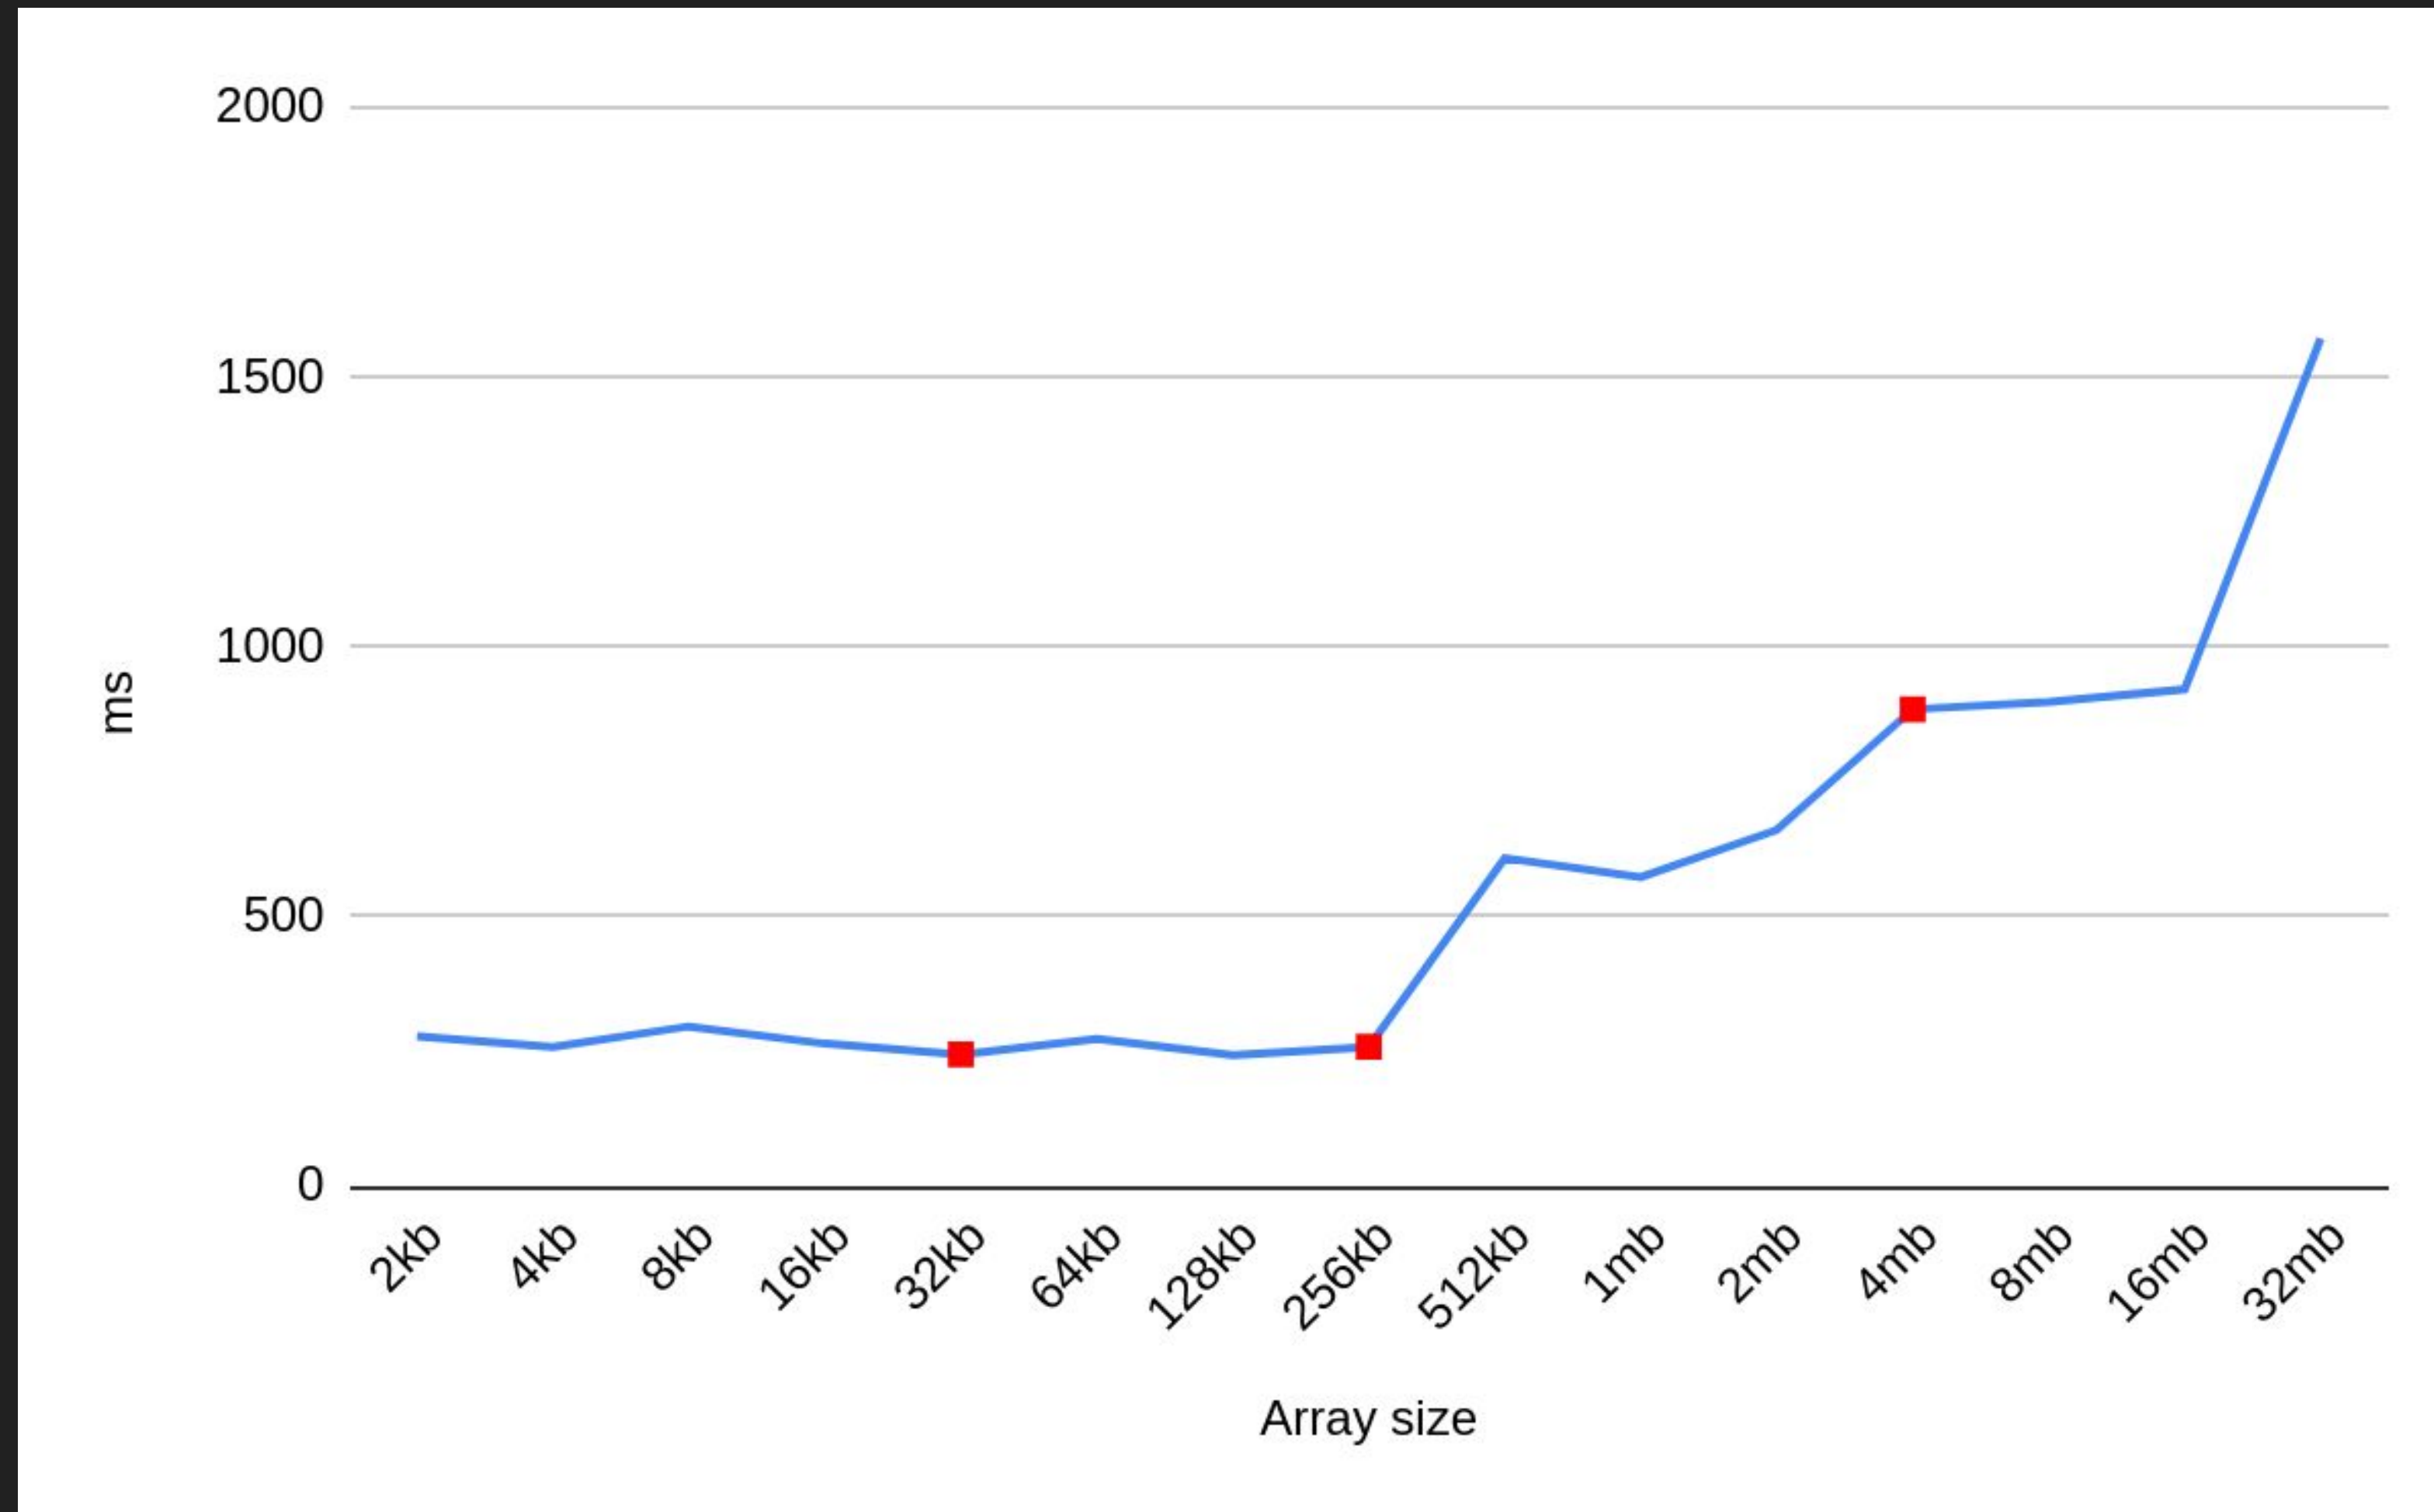
Caches on my Thinkpad

```
> getconf -a | grep CACHE
LEVEL1_ICACHE_SIZE          32768
LEVEL1_ICACHE_ASSOC         8
LEVEL1_ICACHE_LINESIZE     64
LEVEL1_DCACHE_SIZE          32768
LEVEL1_DCACHE_ASSOC         8
LEVEL1_DCACHE_LINESIZE     64
LEVEL2_CACHE_SIZE           262144
LEVEL2_CACHE_ASSOC          4
LEVEL2_CACHE_LINESIZE      64
LEVEL3_CACHE_SIZE           4194304
LEVEL3_CACHE_ASSOC          16
LEVEL3_CACHE_LINESIZE      64
LEVEL4_CACHE_SIZE           0
LEVEL4_CACHE_ASSOC          0
LEVEL4_CACHE_LINESIZE      0
```


Verifying Specs

```
let steps = 256 * 1024 * 2048;  
let length_mod = arr.len() - 1;  
for i in 0..steps {  
    arr[(i * 16) & length_mod] += 1;  
}
```

Verifying Specs



Step Sizes Example

Accessing every element

```
for i in 0..arr.len() {  
    arr[i] *= 3;  
}
```

Accessing every 16th element

```
for i in (0..arr.len()).step_by(16) {  
    arr[i] *= 3;  
}
```

Cachelines



Rows vs Columns

```
for i in 0..A[0].len() {  
    for j in 0..A.len() {  
        count += A[i][j];  
    }  
}
```

VS

```
for i in 0..A[0].len() {  
    for j in 0..A.len() {  
        count += A[j][i];  
    }  
}
```

Rows vs Columns

```
for i in 0..A[0].len() {  
    for j in 0..A.len() {  
        count += A[i][j];  
    }  
}
```

VS

```
for i in 0..A[0].len() {  
    for j in 0..A.len() {  
        count += A[j][i];  
    }  
}
```

Benchmark #1: ./target/release/row

Time (mean \pm σ): 284.5 ms \pm 25.2 ms

Range (min ... max): 263.7 ms ... 327.4 ms

[User: 133.0 ms, System: 150.4 ms]

10 runs

Benchmark #1: ./target/release/column

Time (mean \pm σ): 925.8 ms \pm 62.3 ms

Range (min ... max): 872.5 ms ... 1019.6 ms

[User: 795.6 ms, System: 128.5 ms]

10 runs

More examples

```
for _i in 0..steps {  
    arr[0] += 1;  
    arr[0] += 1;  
}
```

VS

```
for _i in 0..steps {  
    arr[0] += 1;  
    arr[1] += 1;  
}
```

Parallel Instructions

```
for _i in 0..steps {  
    arr[0] += 1;  
    arr[0] += 1;  
}
```

VS


```
for _i in 0..steps {  
    arr[0] += 1;  
    arr[1] += 1;  
}
```

Code snippet on the right is ~2x faster

Yet Another Example

```
let threads: Vec<_> = (0..*P)
    .collect::
```

Yet Another Example




Yet Another Example

```
let threads: Vec<_> = (0..*P)
    .collect::
```


Yet Another Example

Speedup with mutex

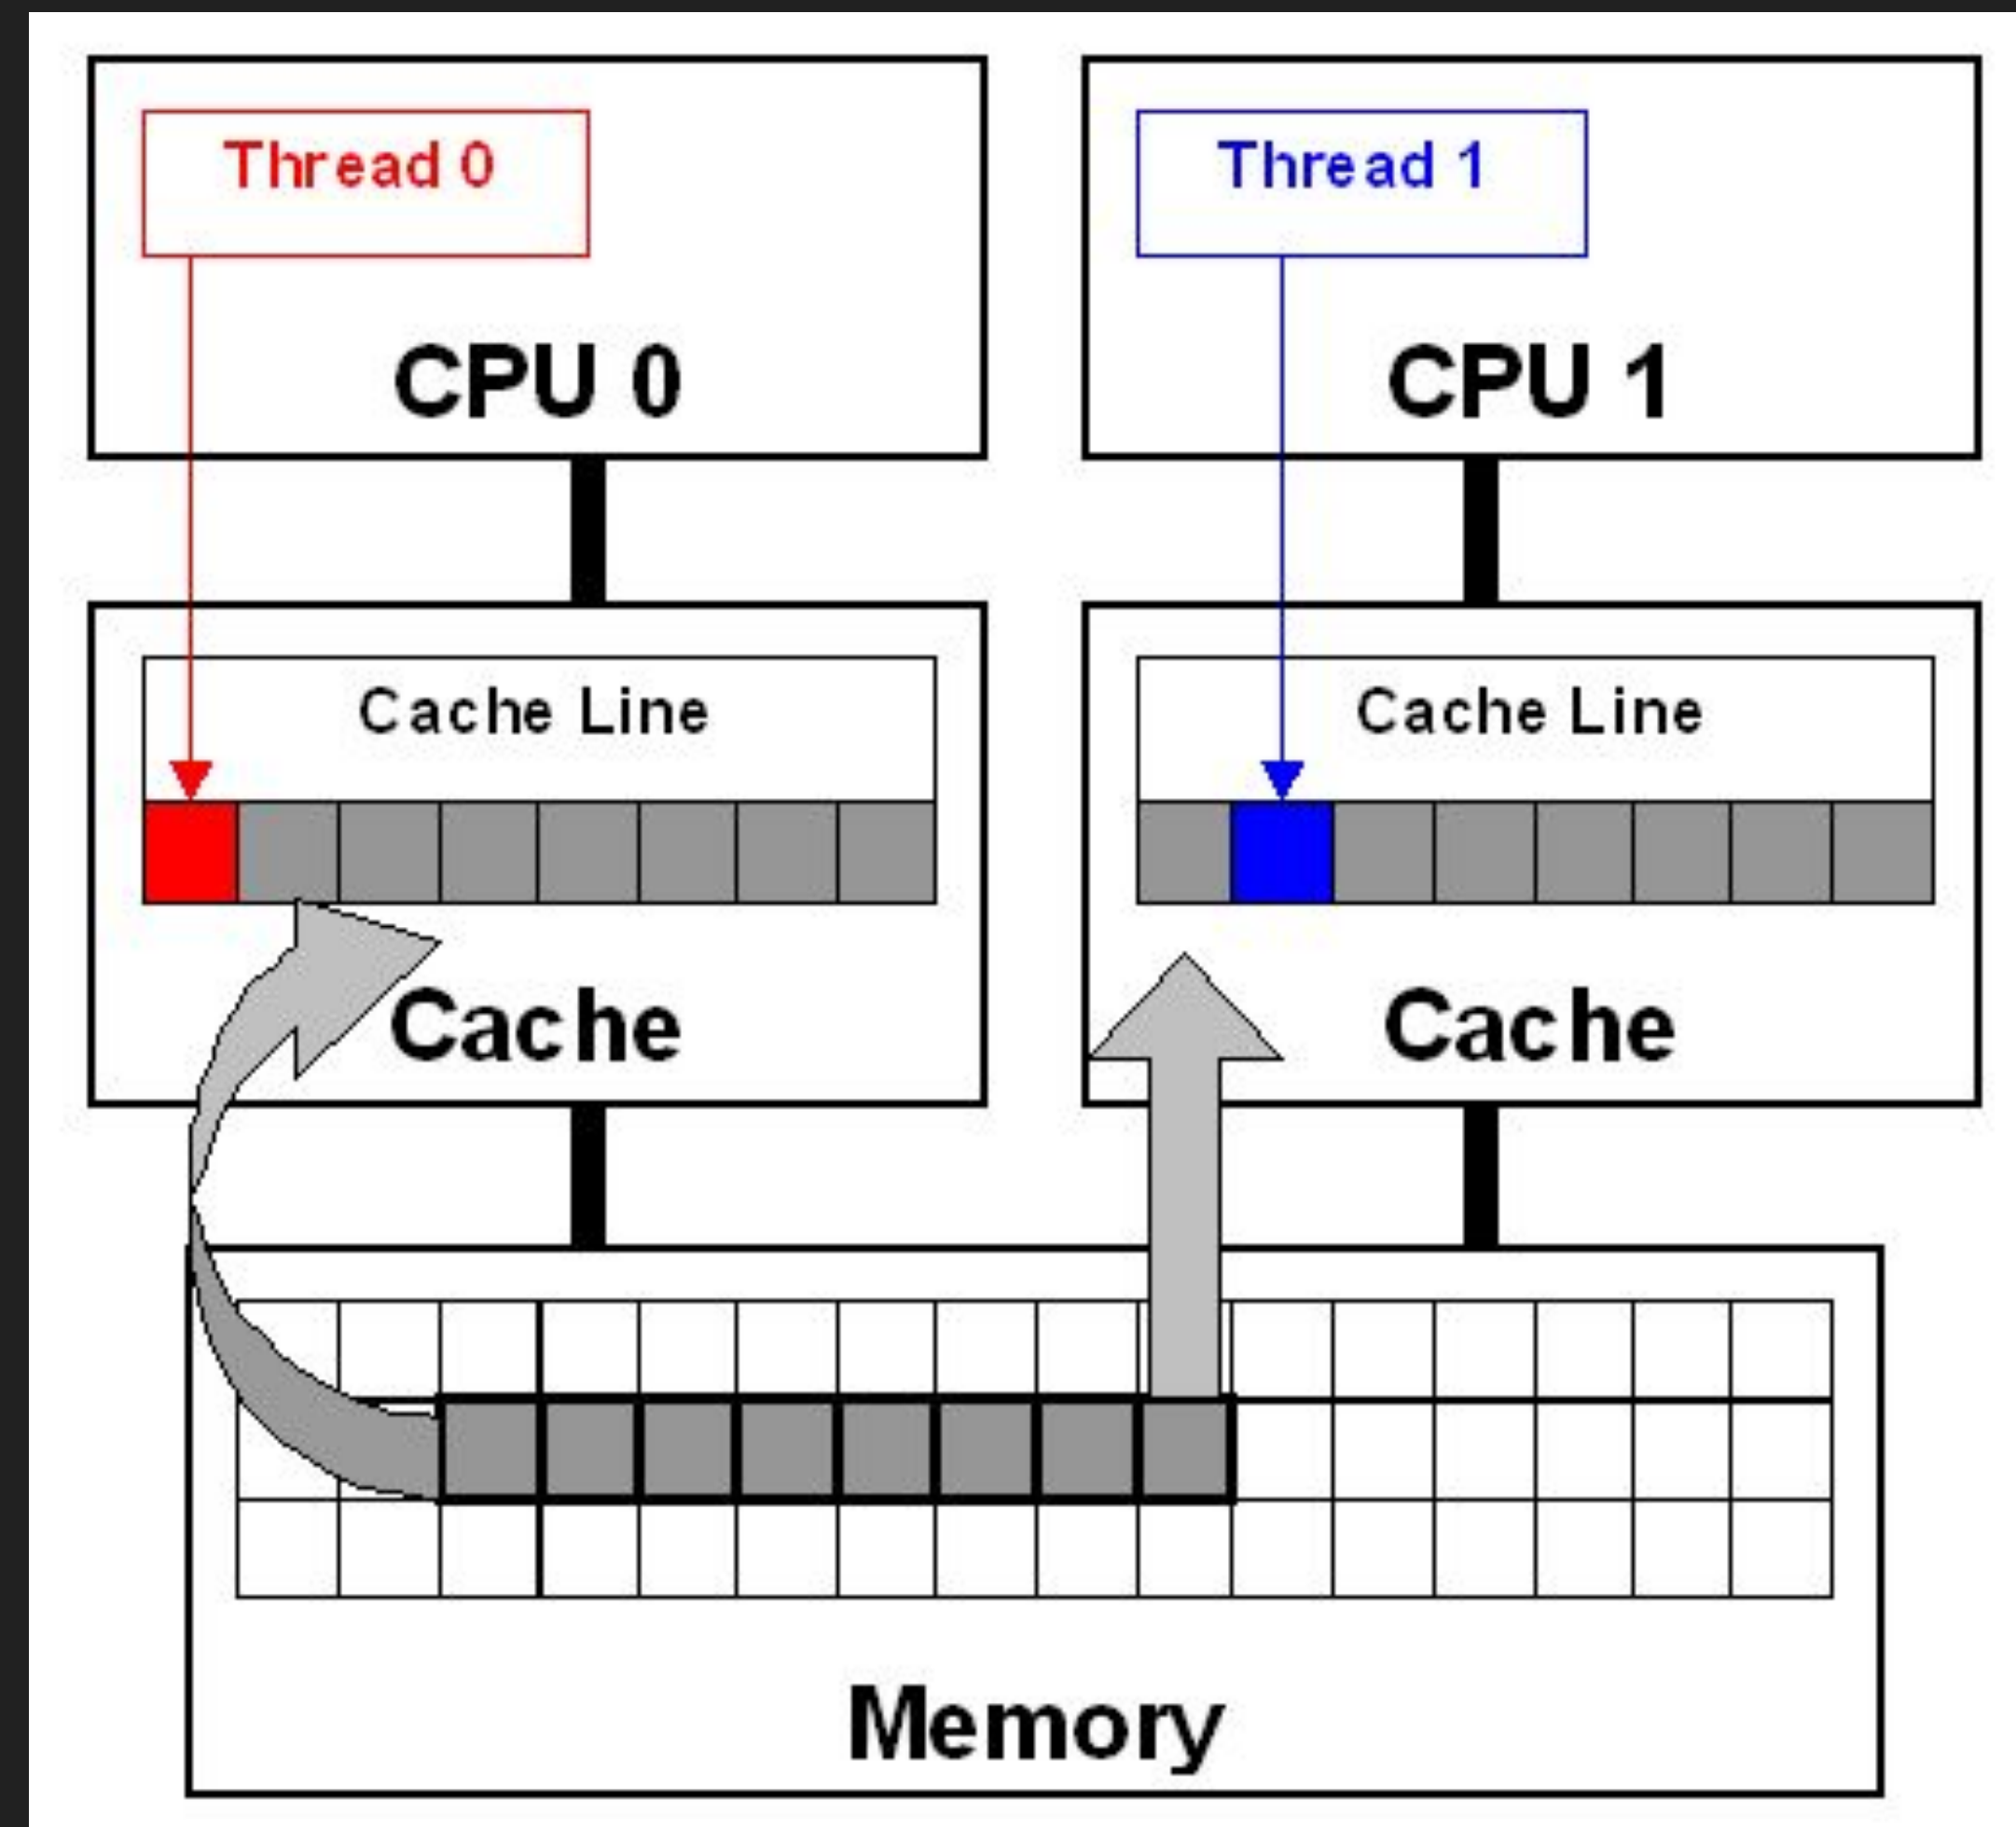


Yet Another Example

Lock-free



False Sharing



Cache Associativity (L1)


- Number of slots
 - L1: 512
- 64-byte chunks partitioned into sets
 - lowest 6 bits determine to which of the 64 sets the cacheline belongs to
 - cache can hold at most 8 lines per any given set

Cache Associativity


```
let rep = 1024 * 1024;
let mut p = 0;
for _i in 0..rep {
    arr[p] += 1;
    p += K;
    if p >= arr.len() {
        p = 0;
    }
}
```


Cache Associativity 8mb Array

Perfomance difference with step size K



Cache Associativity



Conclusions

- Performance behaviour can feel counterintuitive

Conclusions

- Performance behaviour can feel counterintuitive
- Trust the computer but verify

Conclusions

- Performance behaviour can feel counterintuitive
- Trust the computer but verify
- Devs can help computers do a better job

Sources

- <https://akkadia.org/drepper/cpumemory.pdf>
- <http://igoro.com/archive/gallery-of-processor-cache-effects/>
- Scott Meyers: Cpu Caches and Why You Care:
<https://www.youtube.com/watch?v=WDIkqP4JbkE>
- <https://software.intel.com/content/www/us/en/develop/articles/avoiding-and-identifying-false-sharing-a-mong-threads.html>

Questions?

Tommi Jalkanen

github: [koura](#)

twitter: [@om_nommy](#)